

Exercise


Namit Agrawal, Timothy Cheng

08/16/20

Note: Much of the analysis below was inspired by various R files written by Professor James Scott.

Problem 1: Visual story telling: green buildings

It does make sense to use the median rather than mean as the non-green buildings have many outliers as suggested by the boxplot below



The stats guru is right about removing buildings with less than 10% occupancy as based on the summary below. Within the group of buildings that have only 10% occupancy, there is only one building with a green rating. In addition, roughly half of the buildings have 3 stories and very few buildings have a Class A designation. Hence we should remove buildings with less than 10% occupancy as it may distort the analysis.

##	CS_PropertyID	cluster	size	empl_gr	
##	Min. :	57	Min. : 8.0	Min. : 1624	Min. : -1.490
##	1st Qu.: 239800	1st Qu.: 256.0	1st Qu.: 11661	1st Qu.: 1.740	
##	Median : 393481	Median : 371.0	Median : 40000	Median : 2.300	
##	Mean : 1084028	Mean : 536.2	Mean : 62209	Mean : 3.884	

```
## 3rd Qu.: 819966 3rd Qu.:1002.0 3rd Qu.: 83770 3rd Qu.: 3.700
## Max. :6008486 Max. :1230.0 Max. :427383 Max. :67.780
## NA's :1
## Rent leasing_rate stories age renovated
## Min. : 7.00 Min. :0.000 Min. : 1.000 Min. : 0.00 0:148
## 1st Qu.: 16.23 1st Qu.:0.000 1st Qu.: 2.000 1st Qu.: 28.00 1: 67
## Median : 20.50 Median :0.000 Median : 3.000 Median : 57.00
## Mean : 22.44 Mean :1.280 Mean : 4.819 Mean : 54.42
## 3rd Qu.: 27.00 3rd Qu.:0.375 3rd Qu.: 6.000 3rd Qu.: 85.00
## Max. :111.11 Max. :9.780 Max. :19.000 Max. :118.00
##
## class_a class_b LEED Energystar green_rating
## 0:193 Min. :0.0000 Min. :0 Min. :0.000000 0:214
## 1: 22 1st Qu.:0.0000 1st Qu.:0 1st Qu.:0.000000 1: 1
## Median :0.0000 Median :0 Median :0.000000
## Mean :0.4884 Mean :0 Mean :0.004651
## 3rd Qu.:1.0000 3rd Qu.:0 3rd Qu.:0.000000
## Max. :1.0000 Max. :0 Max. :1.000000
##
## net amenities cd_total_07 hd_total07 total_dd_07
## Min. :0.000000 0:189 Min. : 130 Min. : 0 Min. :2103
## 1st Qu.:0.000000 1: 26 1st Qu.: 684 1st Qu.:1419 1st Qu.:2869
## Median :0.000000 Median :1113 Median :2472 Median :4854
## Mean :0.004651 Mean :1676 Mean :3141 Mean :4816
## 3rd Qu.:0.000000 3rd Qu.:2746 3rd Qu.:4916 3rd Qu.:6546
## Max. :1.000000 Max. :5240 Max. :7200 Max. :8244
##
## Precipitation Gas_Costs Electricity_Costs cluster_rent
## Min. :10.46 Min. :0.009487 Min. :0.01782 Min. :10.22
## 1st Qu.:22.71 1st Qu.:0.010118 1st Qu.:0.02453 1st Qu.:18.05
## Median :25.55 Median :0.010296 Median :0.02887 Median :20.74
## Mean :30.34 Mean :0.011579 Mean :0.03111 Mean :23.99
## 3rd Qu.:41.32 3rd Qu.:0.012117 3rd Qu.:0.03781 3rd Qu.:27.02
## Max. :58.02 Max. :0.028914 Max. :0.06278 Max. :65.94
##
```

The median rent for green buildings and non-green buildings is correct if buildings with more than 10% occupancy rate are considered.

```
## 'summarise()' ungrouping output (override with '.groups' argument)
```

```
## # A tibble: 2 x 2
##   green_rating MedianRent
##   <fct>          <dbl>
## 1 0             25.0
## 2 1             27.6
```

Assuming that the building is 250000 square feet, it seems that the stats guru is correct about recuperating the costs in a little under 8 years.

Confounding variables investigation

Renovated Buildings

Based on the numerical summary analysis for buildings that are renovated, it does not seem there is much

confounding going on as the median rents for non renovated and renovated buildings are similar especially for green buildings

```
## 'summarise()' regrouping output by 'renovated' (override with '.groups' argument)
```

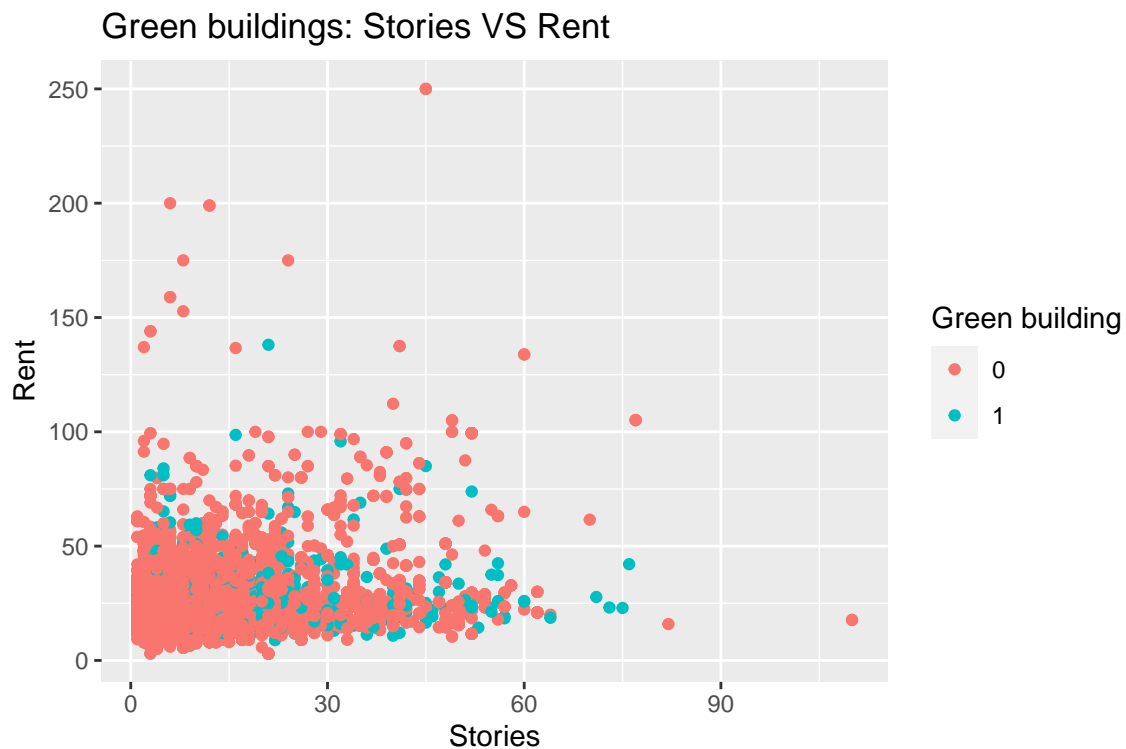
```
## # A tibble: 4 x 3
## # Groups:   renovated [2]
##   renovated green_rating MedianRent
##   <fct>      <fct>          <dbl>
## 1 0        0              27
## 2 0        1             27.6
## 3 1        0             23.5
## 4 1        1             27.0
```

```
## 'summarise()' ungrouping output (override with '.groups' argument)
```

```
## # A tibble: 2 x 2
##   renovated MedianRent
##   <fct>          <dbl>
## 1 0             27
## 2 1             23.8
```

Number of Stories

As suggested by the plot below, the median for stories is a valid selection since there are some outliers. It looks there is not much evidence of confounding for the number of stories in the building, even though there is a slight increase in rent as the number of stories goes up. The median for stories of green buildings only differs by 1, so stories may not directly be affecting the rent.

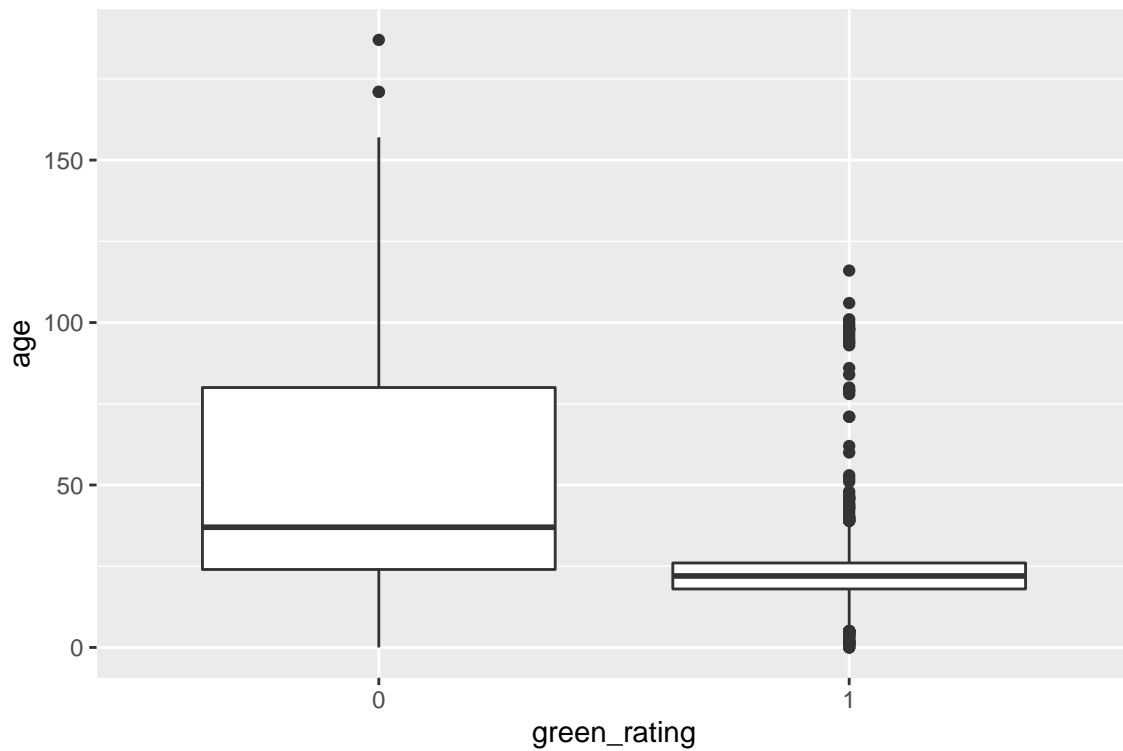


```
## 'summarise()' ungrouping output (override with '.groups' argument)
```

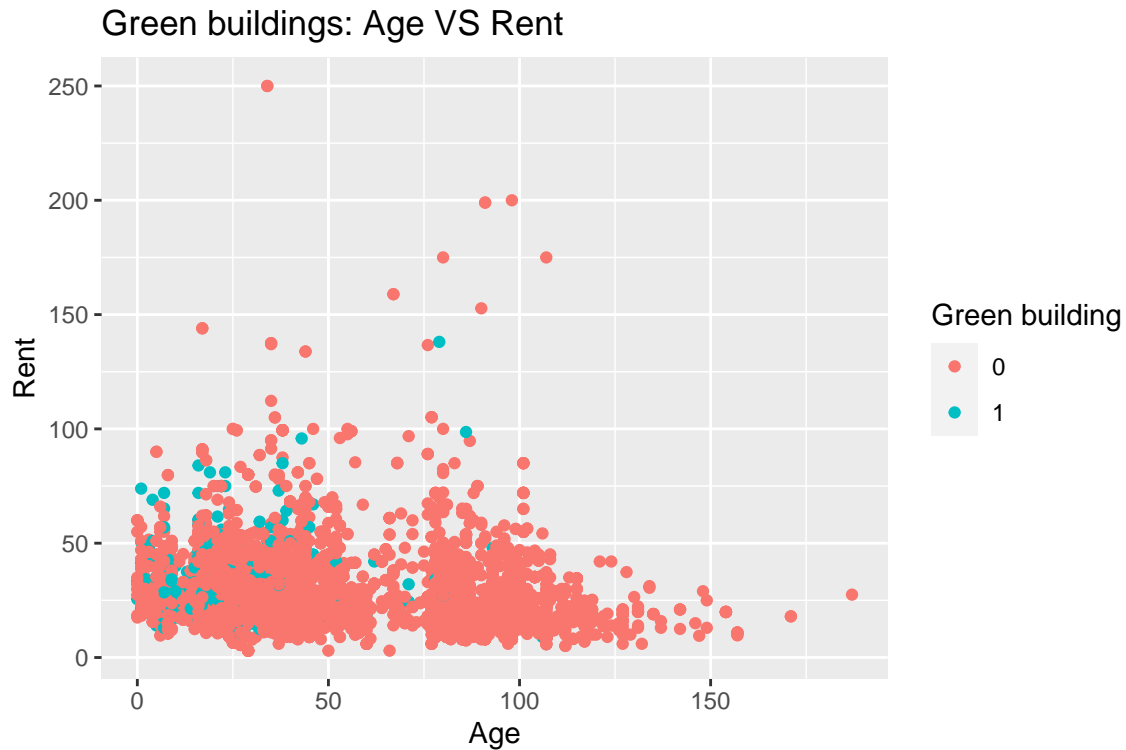
```
## # A tibble: 2 x 2
##   green_rating MedianStories
##   <fct>         <dbl>
## 1 0             10
## 2 1             11
```

Age

An initial analysis provides a stark contrast in age between green and non-green buildings.



However, it looks like there is no confounding for age, as there is no correlation between the age of the building and the rent from the plot below.



```
## 'summarise()' ungrouping output (override with '.groups' argument)
```

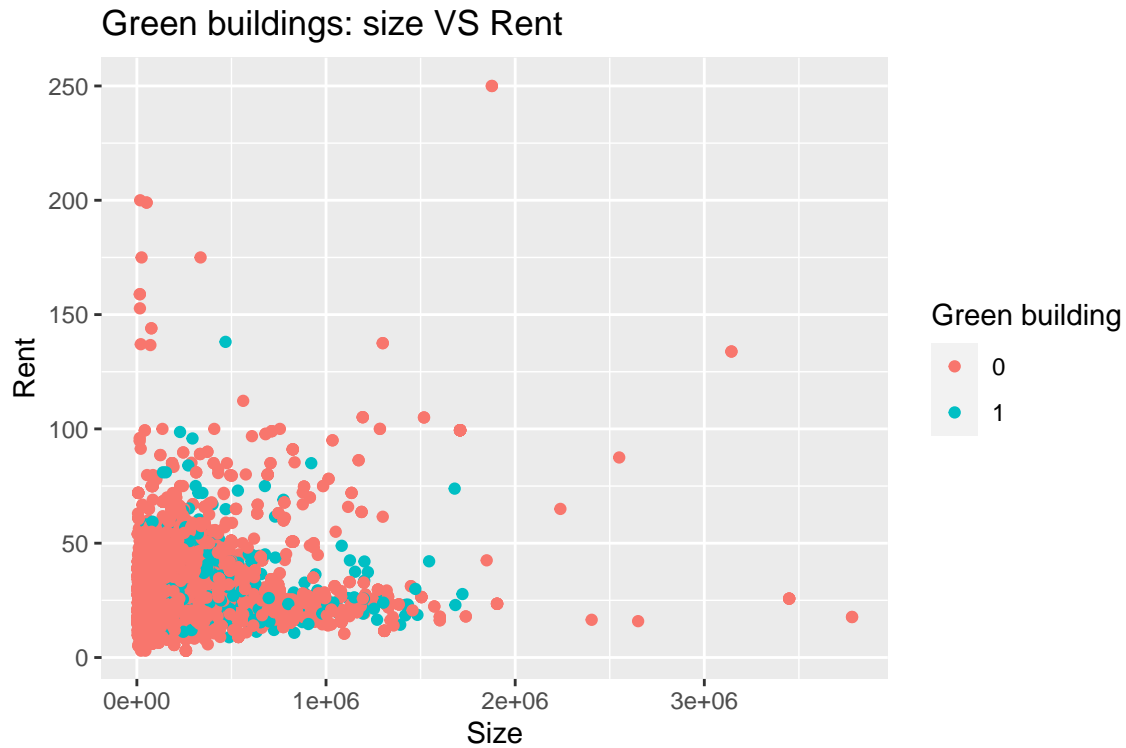
```
## # A tibble: 2 x 2
##   green_rating MedianAge
##   <fct>         <dbl>
## 1 0             36
## 2 1             22
```

Size

It looks like size is definitely a confounding variable, as size is correlated with rent from plot below and the median size for green buildings is double that of non-green. Thus, there is a premium in rent for larger sizes, as expected.

```
## 'summarise()' ungrouping output (override with '.groups' argument)
```

```
## # A tibble: 2 x 2
##   green_rating MedianSize
##   <fct>         <dbl>
## 1 0          123250
## 2 1          241199
```

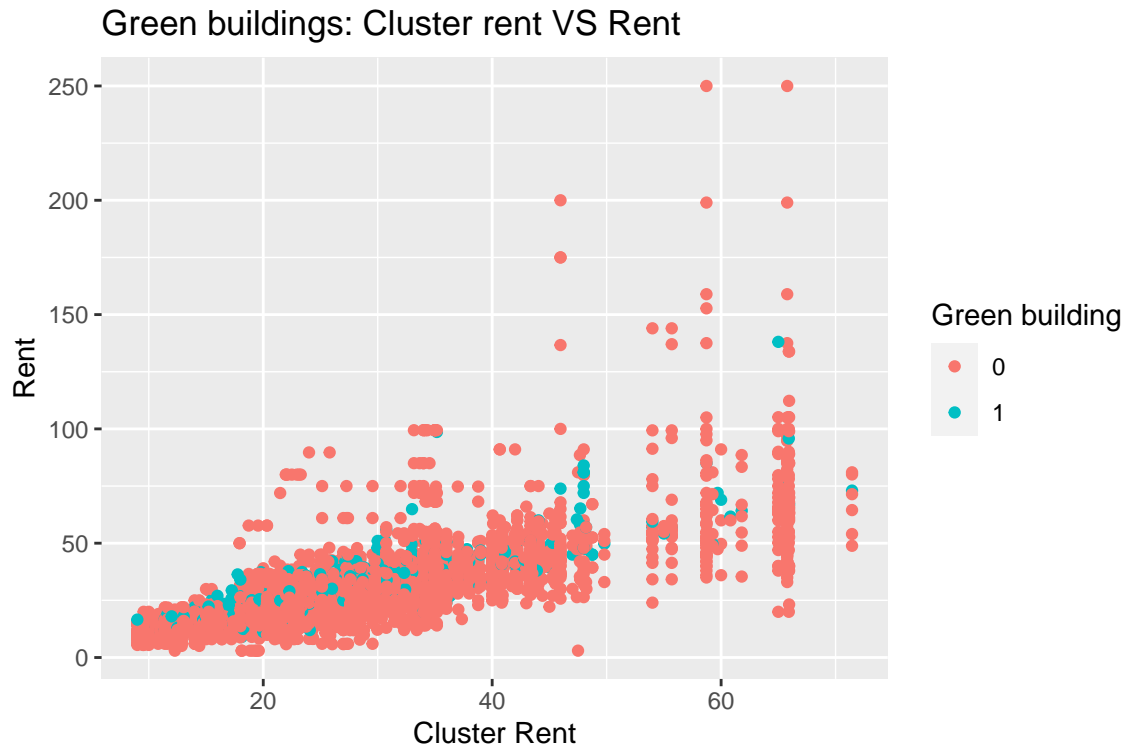


Cluster Rent

There does not seem to be confounding for cluster rent, as the median for cluster rent is approximately the same between green and non-green buildings. However, cluster rent is highly correlated with the rent of the building.

```
## 'summarise()' ungrouping output (override with '.groups' argument)
```

```
## # A tibble: 2 x 2
##   green_rating MedianClusterRent
##   <fct>          <dbl>
## 1 0              25.2
## 2 1              25.4
```

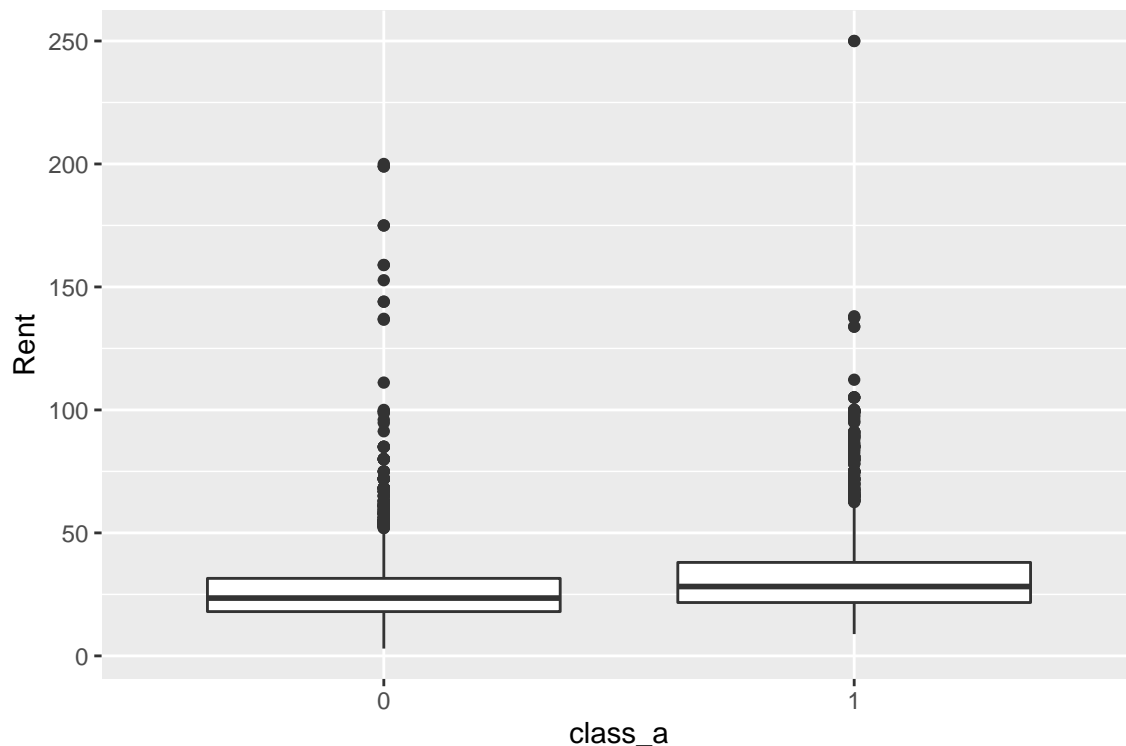


Class

It looks like the Class designation of buildings is a confounding variable, as Class A buildings have generally higher median rents - in addition, having a green_rating with the Class A designation drives median rents even higher. Class A designated buildings seem to correlate with overall rent as well.

```
## 'summarise()' regrouping output by 'class_a' (override with '.groups' argument)
```

```
## # A tibble: 4 x 3
## # Groups:   class_a [2]
##   class_a green_rating MedianRent
##   <fct>    <fct>         <dbl>
## 1 0        0            23.6
## 2 0        1            25.7
## 3 1        0            28.2
## 4 1        1            28.4
```



```
## 'summarise()' ungrouping output (override with '.groups' argument)
```

```
## # A tibble: 2 x 2
##   class_a MedianRent
##   <fct>      <dbl>
## 1 0          23.9
## 2 1          28.2
```

Thoughts

From the investigation above it seems like size and class are the only confounding variables, as the green buildings tend to have larger spaces and larger spaces have higher rent. However, the stats guru is only taking into account the median rent of all the building with more than 10% occupancy. If we apply another filter to include only 15 story buildings, we see that the rent goes up drastically for green buildings - all the way to 37 dollars. However, it may not be wise to use this filter as there is only 10 green buildings that have 15 stories.

```
## 'summarise()' ungrouping output (override with '.groups' argument)
```

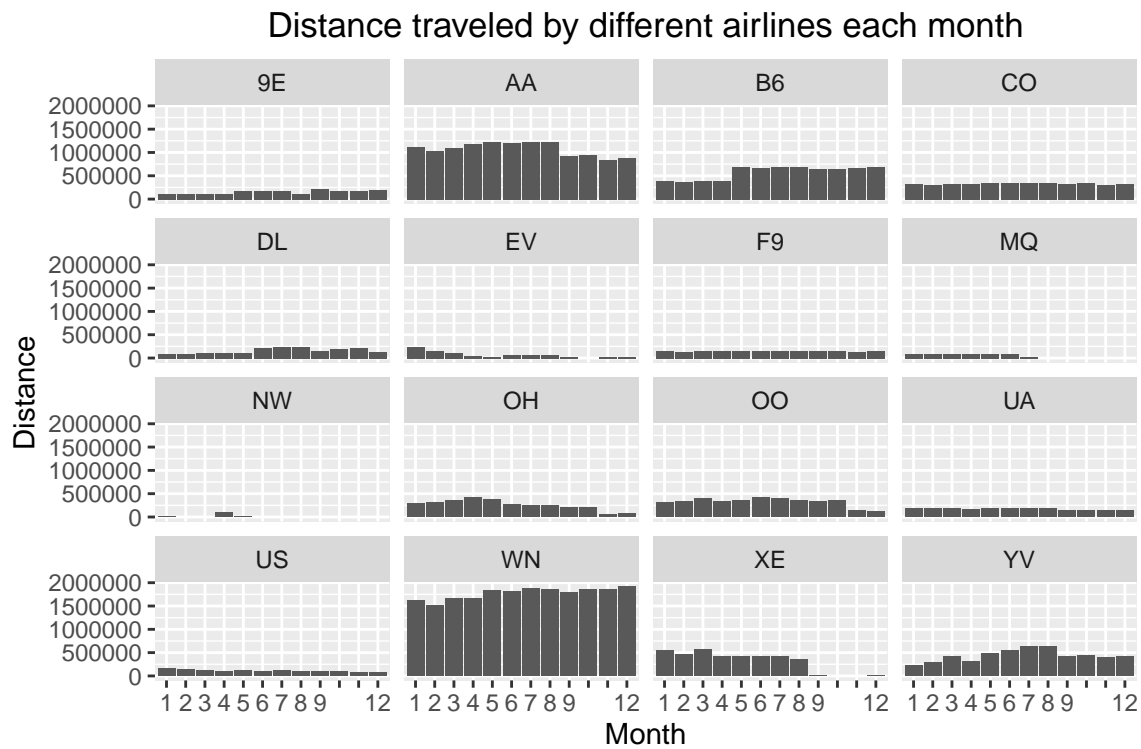
```
## # A tibble: 2 x 3
##   green_rating MedianRent  num
##   <fct>          <dbl> <int>
## 1 0             24.4   156
## 2 1             37.0    10
```

In general, the guru is correct with his analysis, but the analysis is performed on a dataset with a large range of different building specifications. For example, the dataset only contains 10 green buildings that have 15 stories. A larger sample size that adheres to the developers desired specs would provide more valid results.

Problem 2: Visual story telling: flights at ABIA

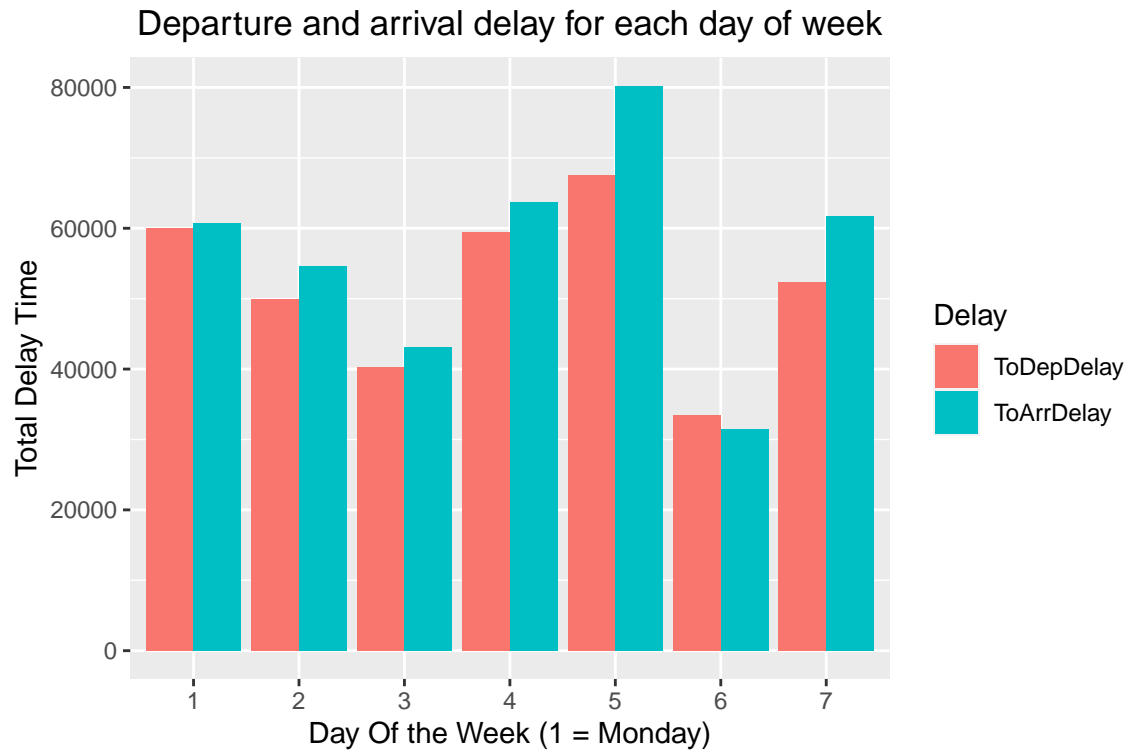
Most active airlines

Here we analyze which airlines are most active throughout the year in terms of the distance flown. As seen by the plot, it appears Southwest (WN) and American Airlines(AA) are the most active in flying out of ABIA and flying to ABIA.



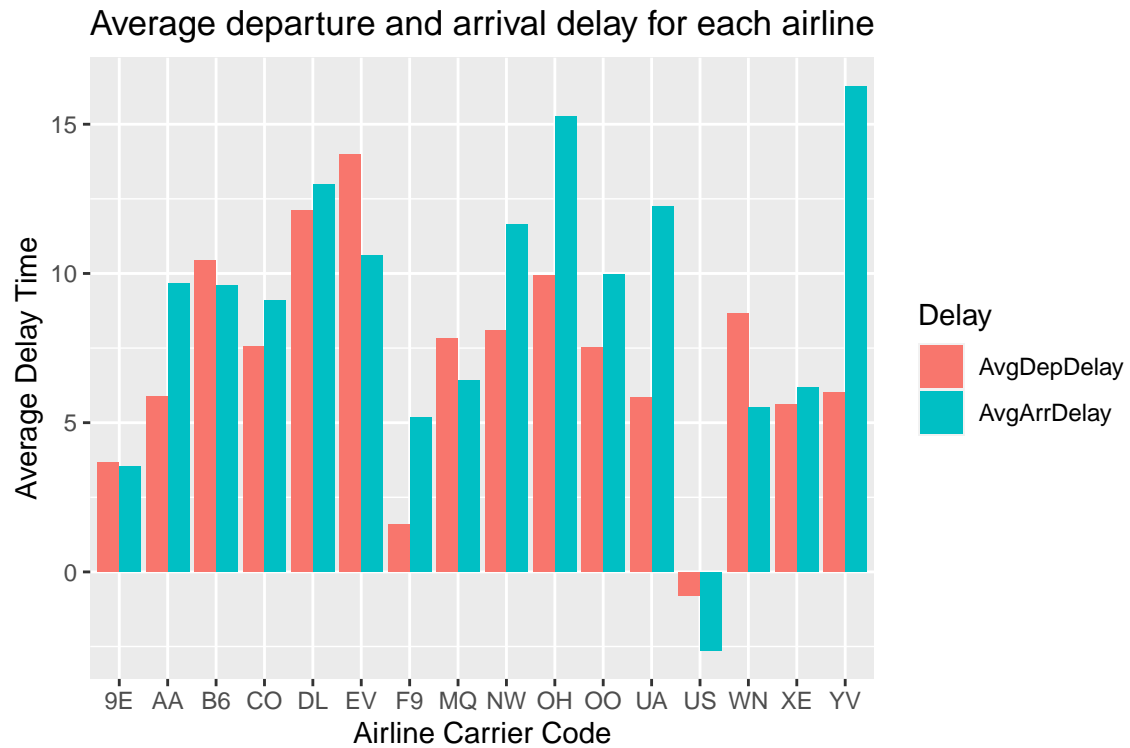
Delays by Day of Week

In the plot below, we analyze the departure and arrival delay for each day of the week. There are more arrival delays than departure delays and Friday is the worst day to travel to/from Austin.



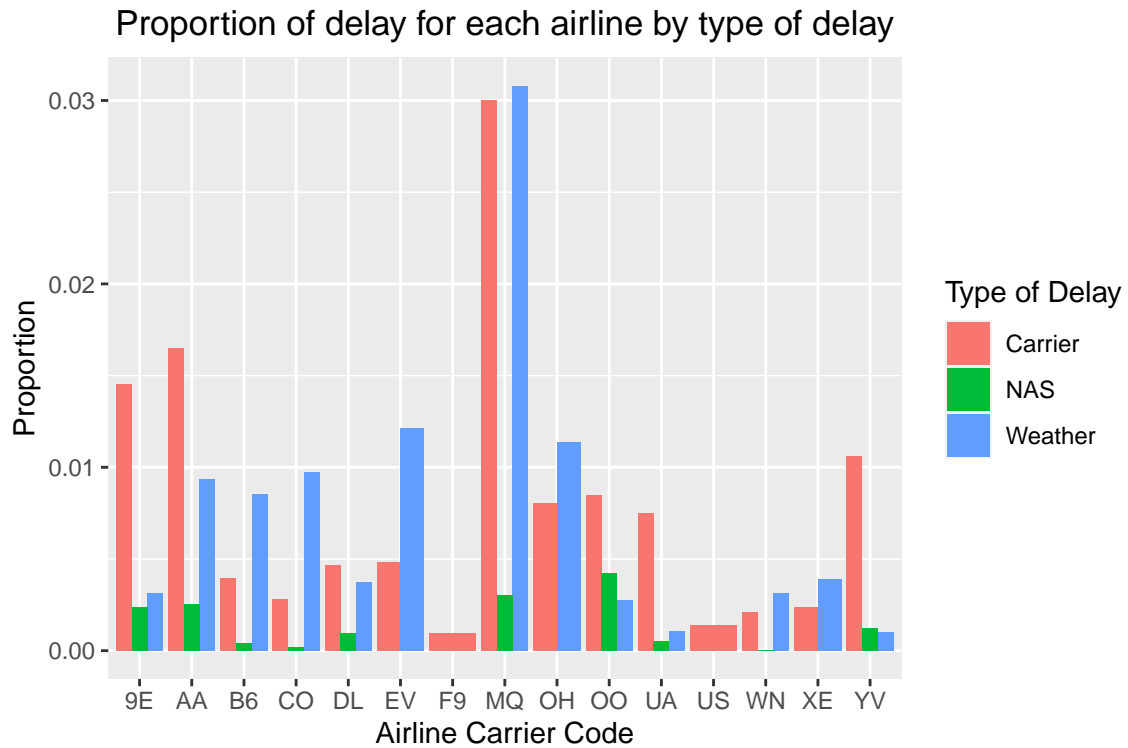
Average Delays per Airline

This plot analyzes the average departure and arrival delay for each airline. It looks like Piedmont Airlines (US) arrives and departs early on average as the delay time is negative.



Most common Delay types by Airline

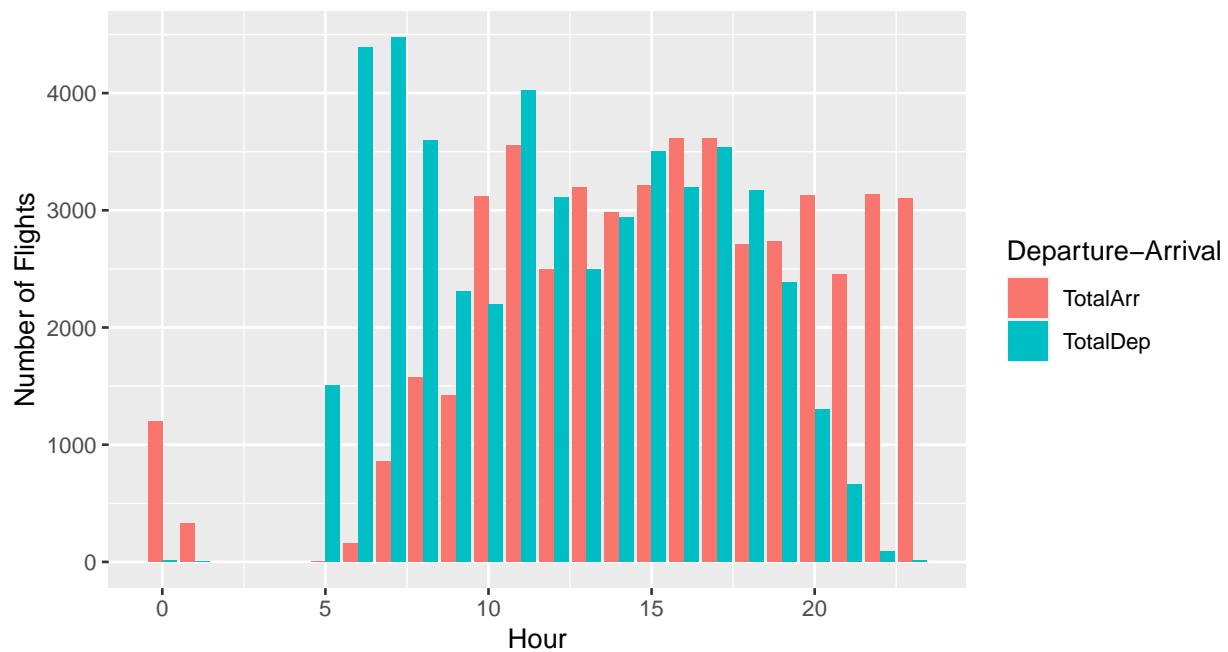
Here we analyze the proportions of delays by airlines via the type of delay. It look like most airlines suffer from carrier and weather delays.



Flights per hour of the Day

The plot below analyzes what are most frequent departure and arrival times. Passengers typically fly *out* early in the morning and fly *in* late at night. Between noon and evening there is an even split between passengers flying in and out.

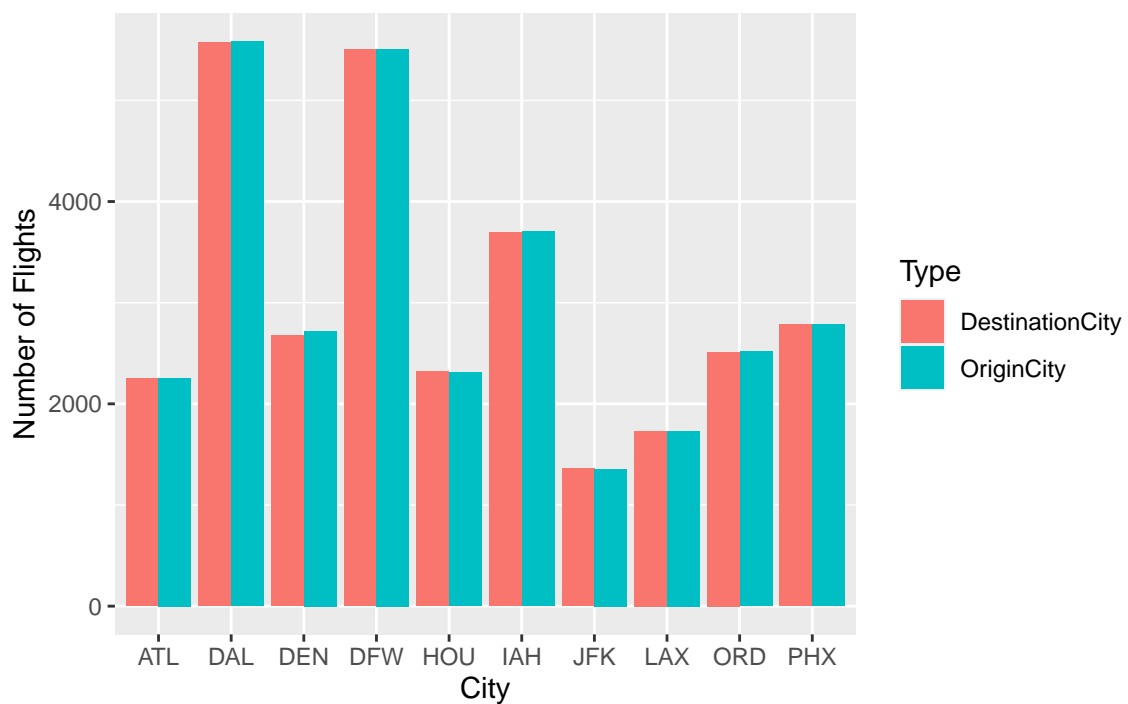
Number of flights at ABIA for each hour by departure and arrival



Flights by City

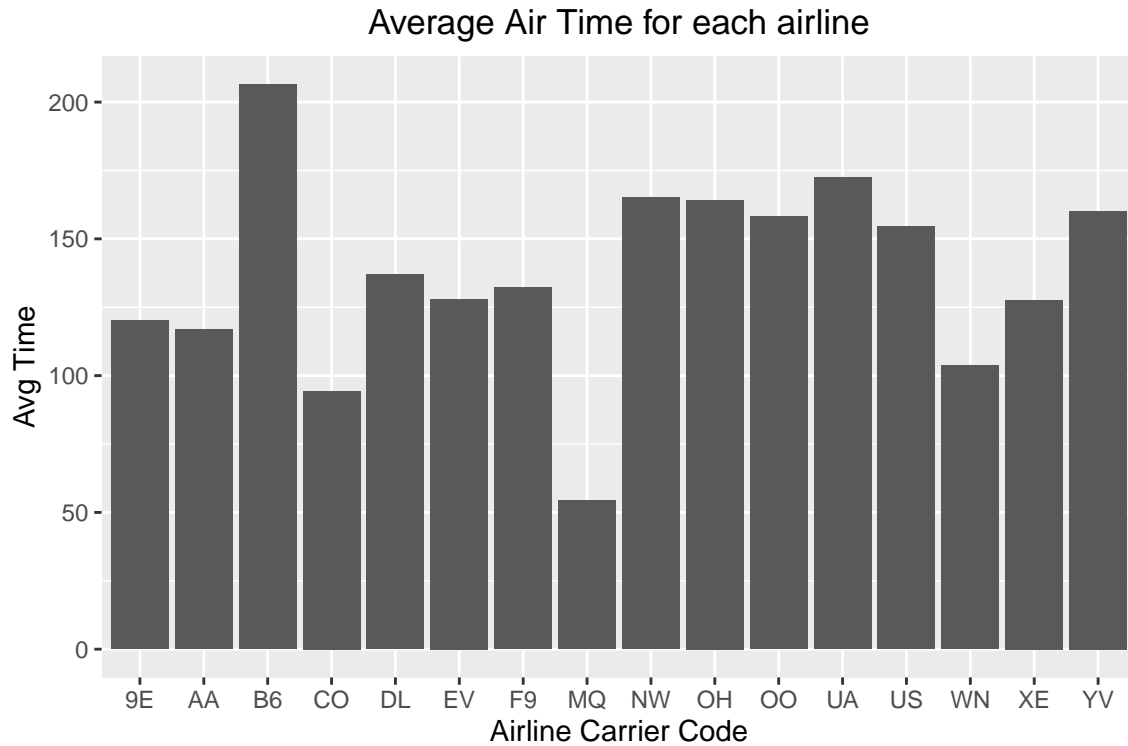
The plot below analyzes the top ten airports to which passengers fly to and fly in from. Dallas and Houston are by far the most popular destinations.

Top 10 cities for incoming and outgoing flights



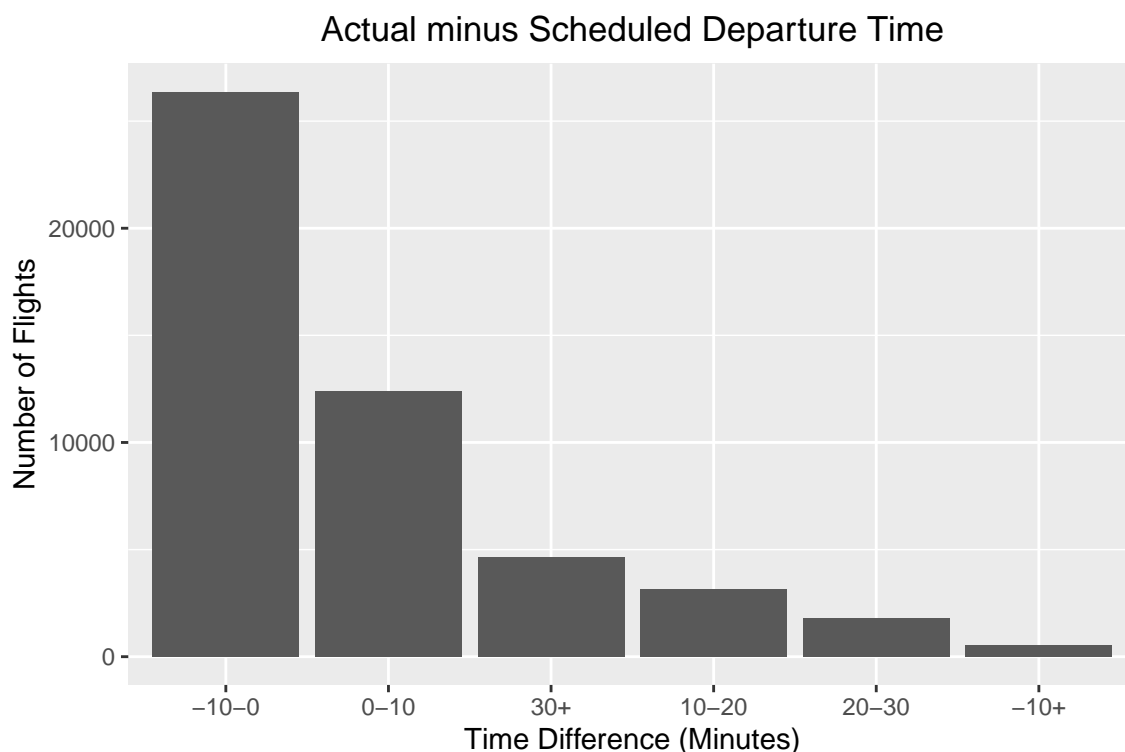
Average time spent flying by Airline

The plot below analyzes the average time spent flying for each airline. JetBlue (B6) flies for more than 3 hours on average.



Actual minus Scheduled Departure Times

The plot below analyzes on average how often an airline deviates from its scheduled departure time. Most airlines leave between 0 to 10 minutes earlier than scheduled!



Problem 3: Portfolio modeling

Background

For this problem, we are analyzing five different ETFs ranging from Gold ETFs to Oil related ETFs. We chose these ETFs because we wanted to explore a diverse set of ETFs based on volatility.

We have chosen to go with 5 ETFs:

“GLD” - The Fund seeks to achieve the performance of gold bullion less the expenses of the Fund

“USO” - The Fund seeks to reflect the performance of the spot price of West Texas Intermediate light, sweet crude oil delivered to Cushing, Oklahoma by investing in a mix of Oil Futures Contracts and Other Oil Interests.

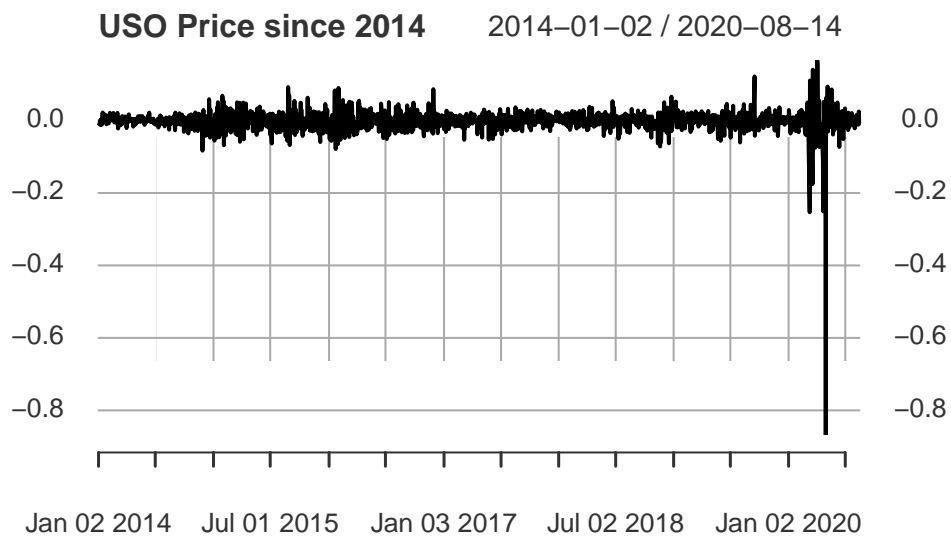
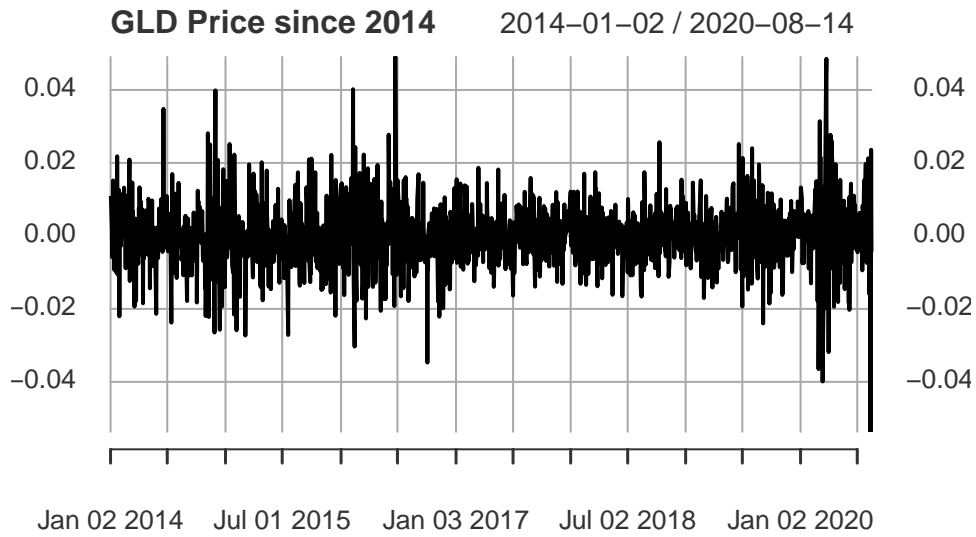
“VNQ” - The Fund seeks to provide a high level of income and moderate long-term capital appreciation by tracking the performance of a benchmark index that measures the performance of publicly traded equity REITs and other real estate-related investments.

“BNO” - BNO tracks the Brent oil spot price using near-month ICE futures contracts.

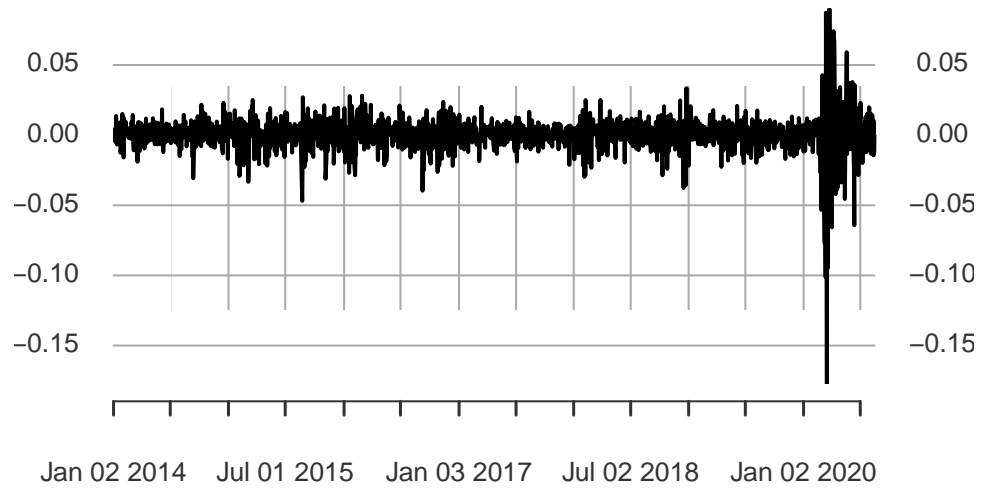
“SLV” - The Fund seeks to reflect generally the performance of the price of silver.

Volatility

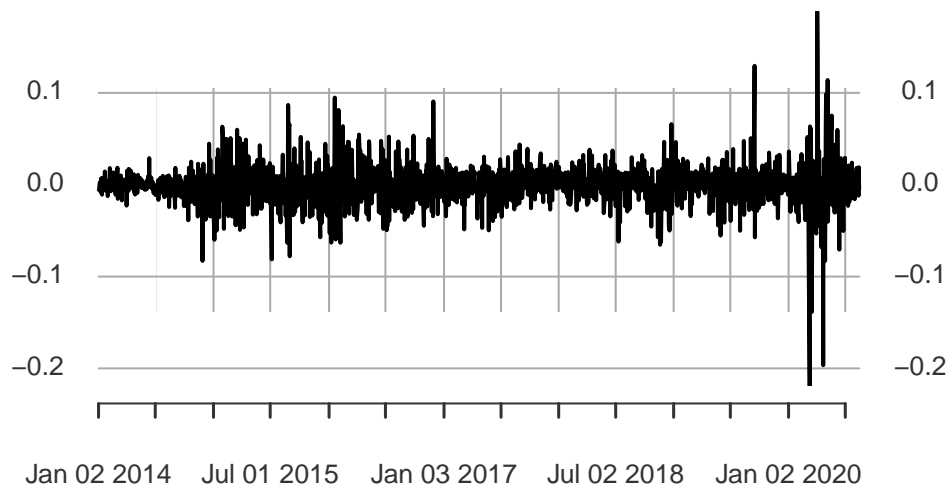
Below are a few plots for the closing prices of ETF. The oil ETFs are the most volatile of the five funds chosen.

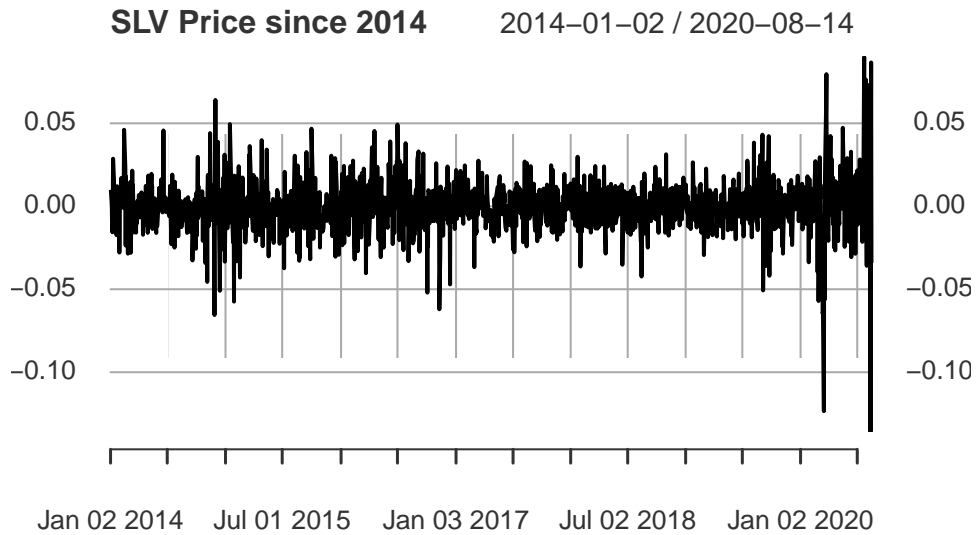


VNQ Price since 2014 2014-01-02 / 2020-08-14




BNO Price since 2014 2014-01-02 / 2020-08-14





Portfolios

Portfolio 1 : A portfolio of equal weights to all ETFs (i.e, 20 percent to all ETFs)

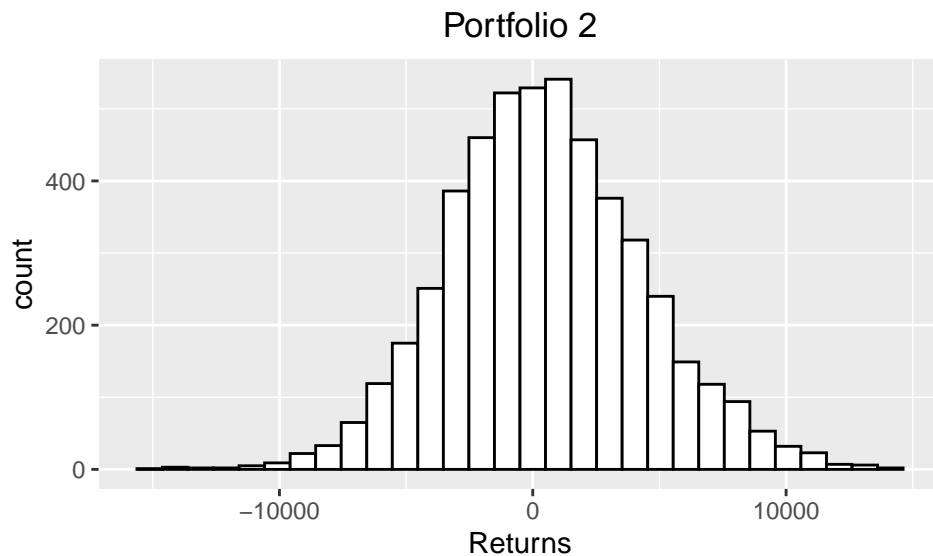


```
##          5%
## 9673.958
```

```
## [1] -391.6459
```

The 5% value at risk for this particular portfolio is roughly \$9,674. The average return on the portfolio is a loss of \$392.

Portfolio 2 : A portfolio that invests 96 percent of wealth into gold and 1 percent into each of the remaining 4 ETFs.

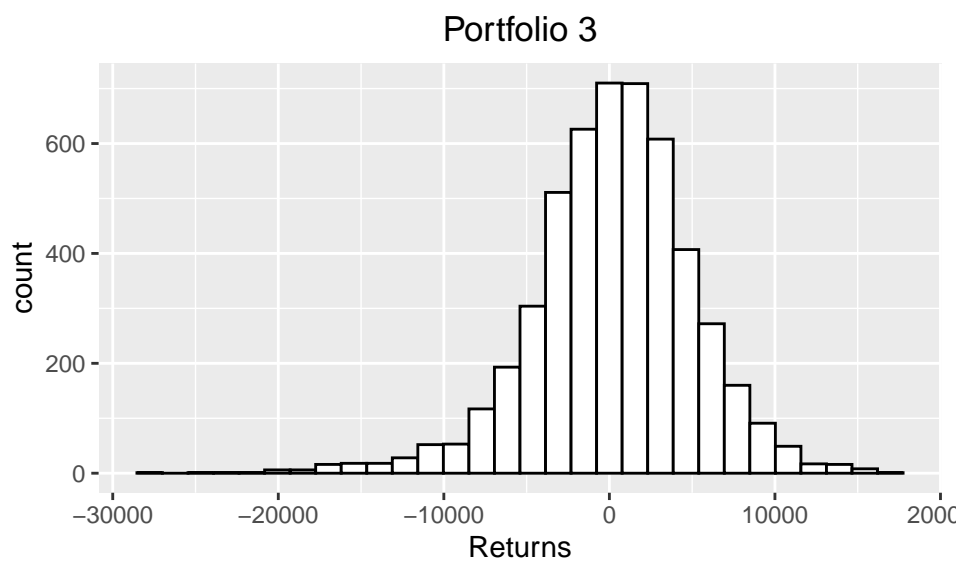


```
##      5%
## 5670.255
```

```
## [1] 490.9026
```

The 5% value at risk for this particular portfolio is roughly \$5,670. The average return on the portfolio is a profit of \$491.

Portfolio 3 : A portfolio that invests 60 percent of wealth into VNI and 10 percent into each of the remaining 4 ETFs.



```
##      5%
## 7694.568
```

[1] 226.436

The 5% value at risk of this particular portfolio is roughly \$7,695. The average return on this portfolio is a profit of \$226.


Report

Based on our analysis, portfolio 2 performed the best. By investing 96% of our wealth into the gold ETF, we were able to achieve the highest returns and the lowest VaR (value at risk) at 5% between all portfolios. This is an interesting result as diversification of the portfolio hurt our investments which suggests that ETFs related to Oil and Silver are significantly more volatile than Gold. This also suggests that Gold is typically a safe investment to make.

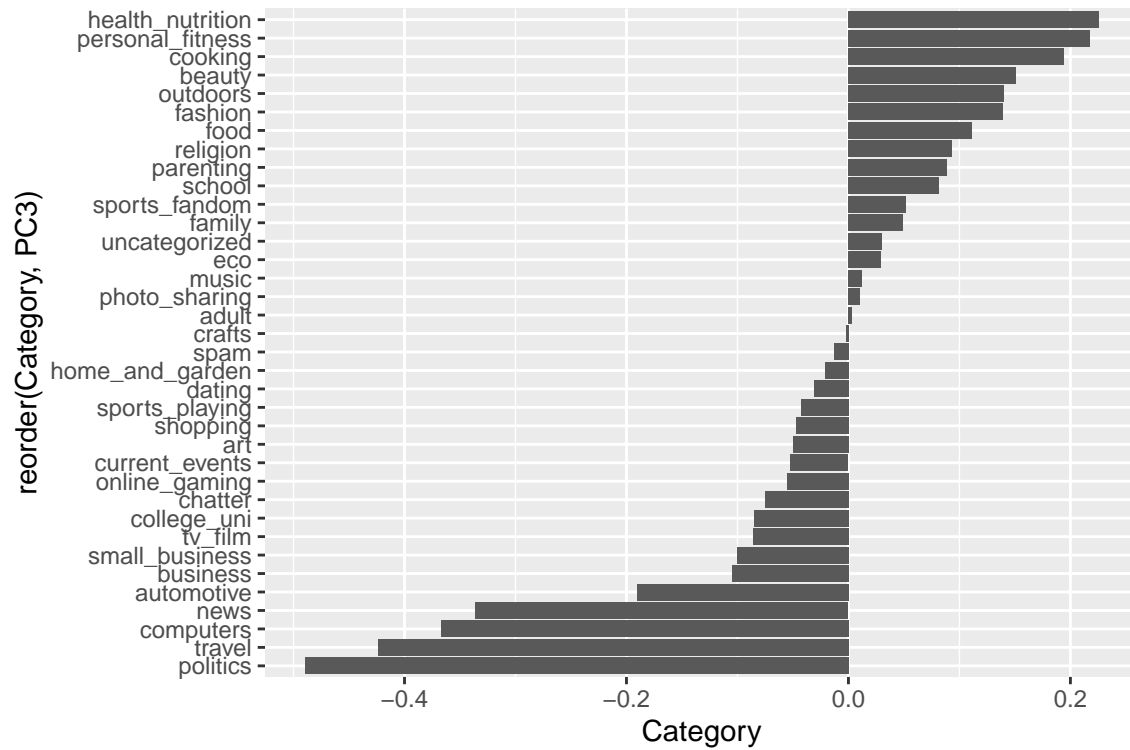
Problem 4: Market Segmentation

From the dataset provided by the company “NutrientH20”, we hope to extract some vital market information regarding the types of followers that “NutrientH20” has.

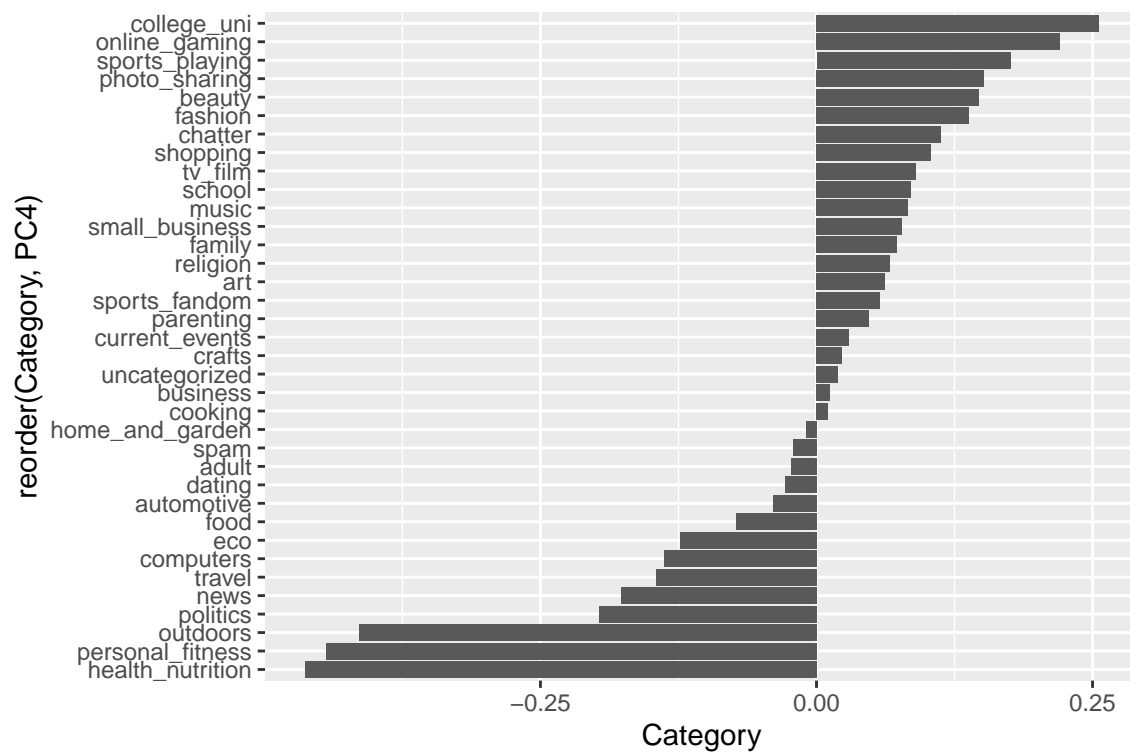
We first perform dimension reduction on the dataset to improve computational ability. In addition, we can visualize the marginal variance explained by adding another PC. Because the elbow is not clear in this plot, we choose a value of 15 PCs to continue our analysis.




Based on our understanding of the dataset, we can conclude which PCs are associated with separating which types of followers. For example, in the plot below, we can see that the third PC has weights that are strongly positive for people that are interested in fitness, but not so much in computers/gaming/politics.



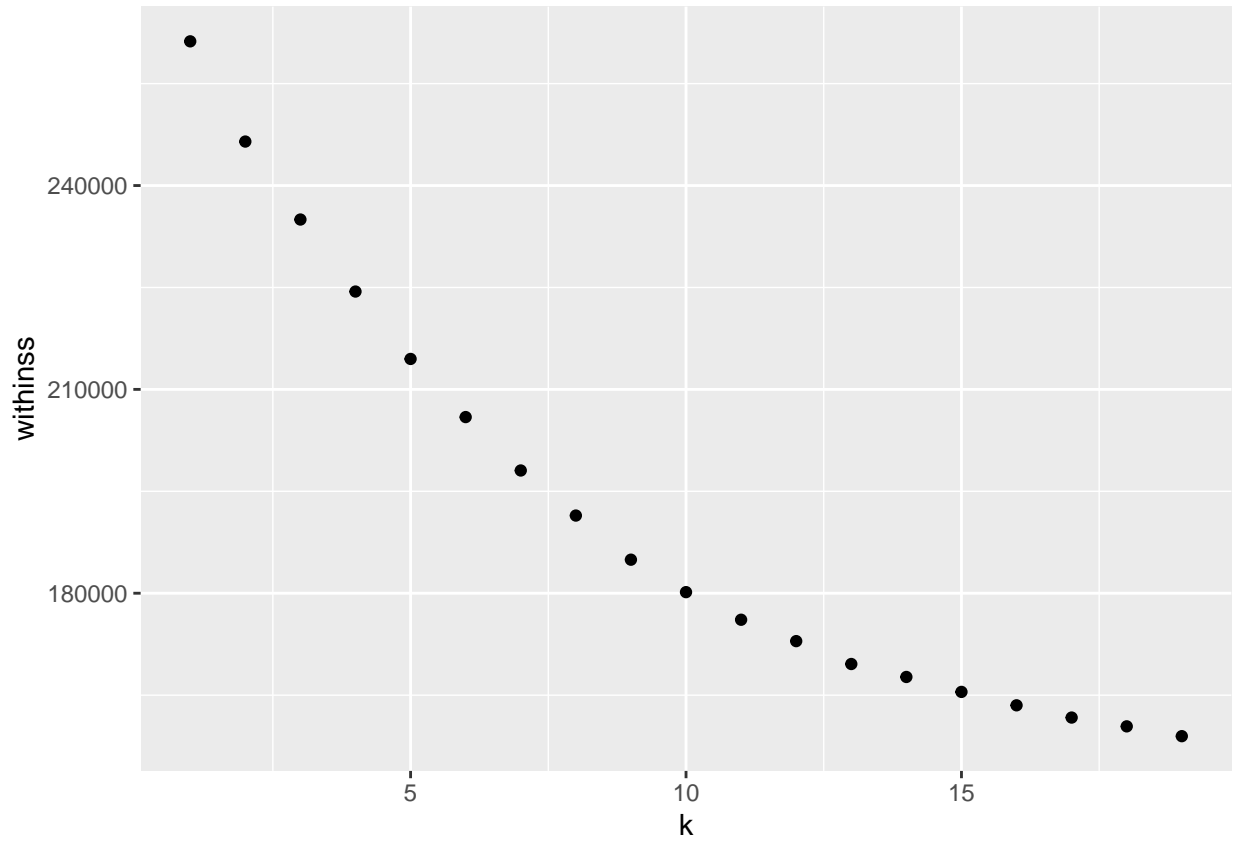
In contrast, our fourth PC (below) has strong negative weights for health and fitness and seems to value online gaming and sports.



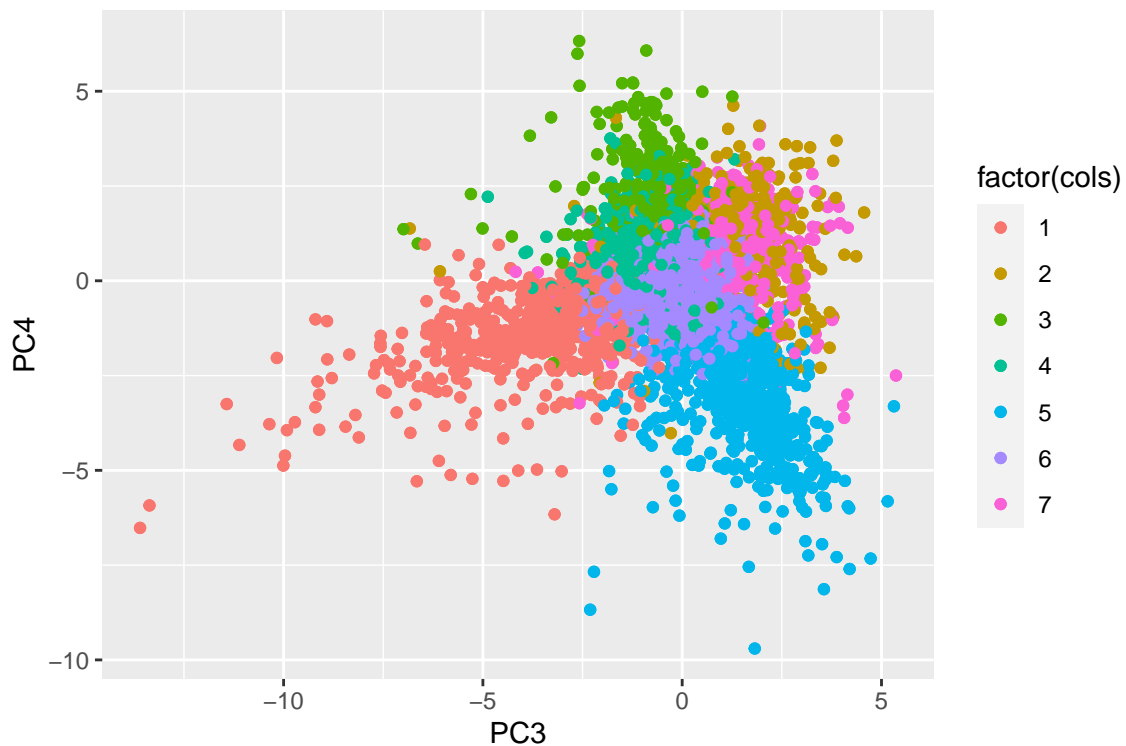
These two PCs, then would be good at separating and visualizing different types of followers in a 2-Dimensional space.



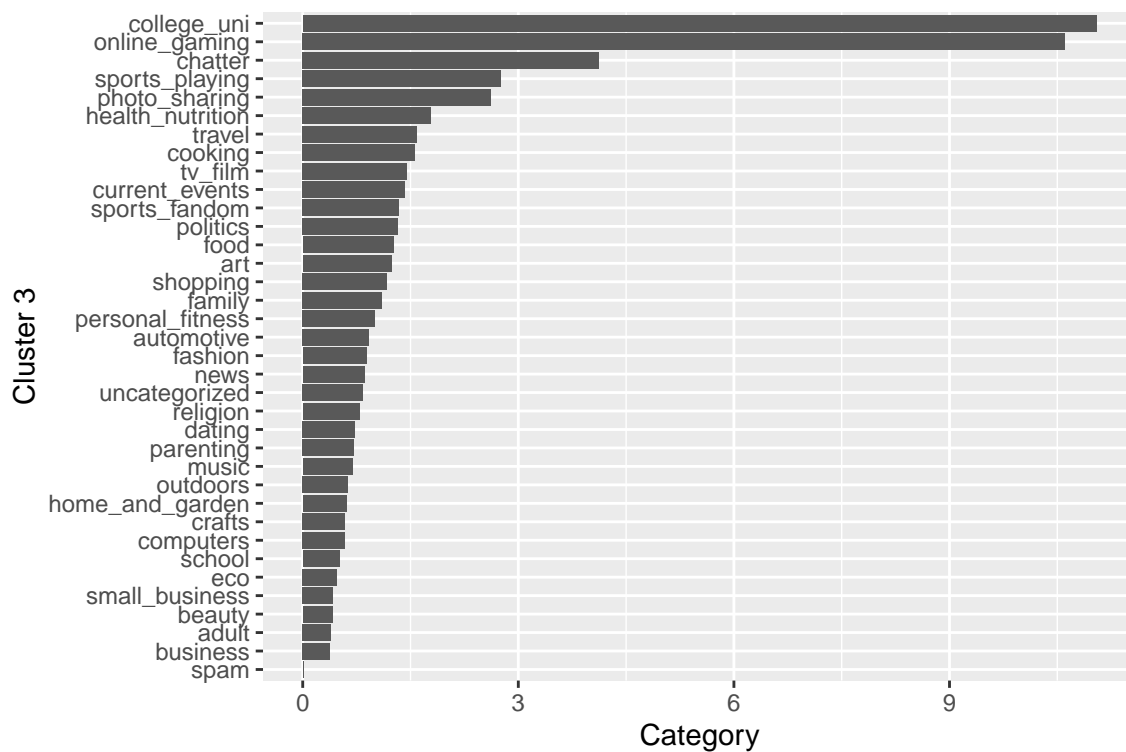
In order to determine clusters of followers in this space, we perform kmeans clustering on our points in PC space. To determine the optimal hyperparameter for clustering, we try several values of k and measure the best one using our within sum of squares as the metric.

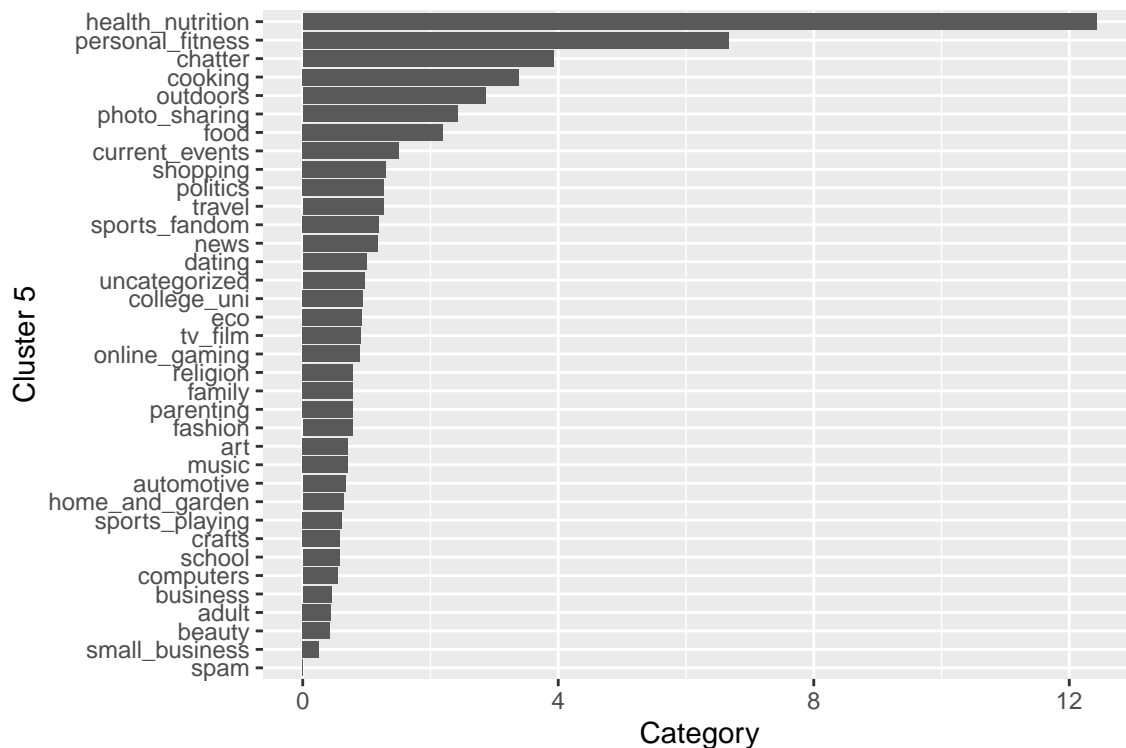


Based on the plot, we decided to choose $k = 7$. We can now visualize the different clusters in the 2D PC space we chose earlier. There seems to be a clear separation between clusters - those with a higher positive value for PC3 but low positive value for PC4 are individuals who love fitness (cluster 5). Those who (we suspect) are college students are going to be in the positive PC4 range with negative PC3 values (cluster 4).



A closer examination of the clusters as determined by kmeans gives us a better picture for what type of followers are within this group.





Problem 5: Author Attribution

Using the ReutersC50 dataset, we hoped to fit models that could predict the author of a particular article based on the article’s textual content alone.

Data Pre-Processing

We first needed to ensure that the test and training datasets had the appropriate format. The datasets need to be formatted into count matrices, where each i, j entry represents the number of j words that the i^{th} document contains (using DocumentTermMatrix function from the “tm” package in R).

Using several pre-processing techniques found in the “tm” package, we convert each word within each article into lower case. We also strip each word of numbers, punctuation, and whitespaces, before then creating the data matrices.

Since both the training and testing data matrices are sparse, we filter out terms that have more than 98% sparse terms (0’s). We choose this filter as it preserves words that specific authors use, which would benefit our predictions. We use the removeSparseTerms function from the “tm” package in R.

Given the filtered dataset, we then convert each i, j entry into a TF-IDF score using the weightTfIdf function in the “tm” package.

In addition, we keep track of which document belongs to which author by creating an additional vector (for both the test and train datasets) where the i^{th} element of the list is the author of document i .

Below is a sample of our training dataset:

```
## <<DocumentTermMatrix (documents: 5, terms: 5)>>
## Non-/sparse entries: 10/15
## Sparsity          : 60%
## Maximal term length: 11
## Weighting         : term frequency - inverse document frequency (normalized) (tf-idf)
## Sample           :
##      Terms
## Docs  access      also  announced authorities  become
##   1 0.03280211 0.006312011 0.02172606 0.03236174 0.02534424
##   2 0.00000000 0.010062627 0.01731788 0.00000000 0.00000000
##   3 0.17601133 0.000000000 0.00000000 0.00000000 0.00000000
##   4 0.00000000 0.000000000 0.01541850 0.00000000 0.00000000
##   5 0.00000000 0.000000000 0.01942981 0.00000000 0.00000000

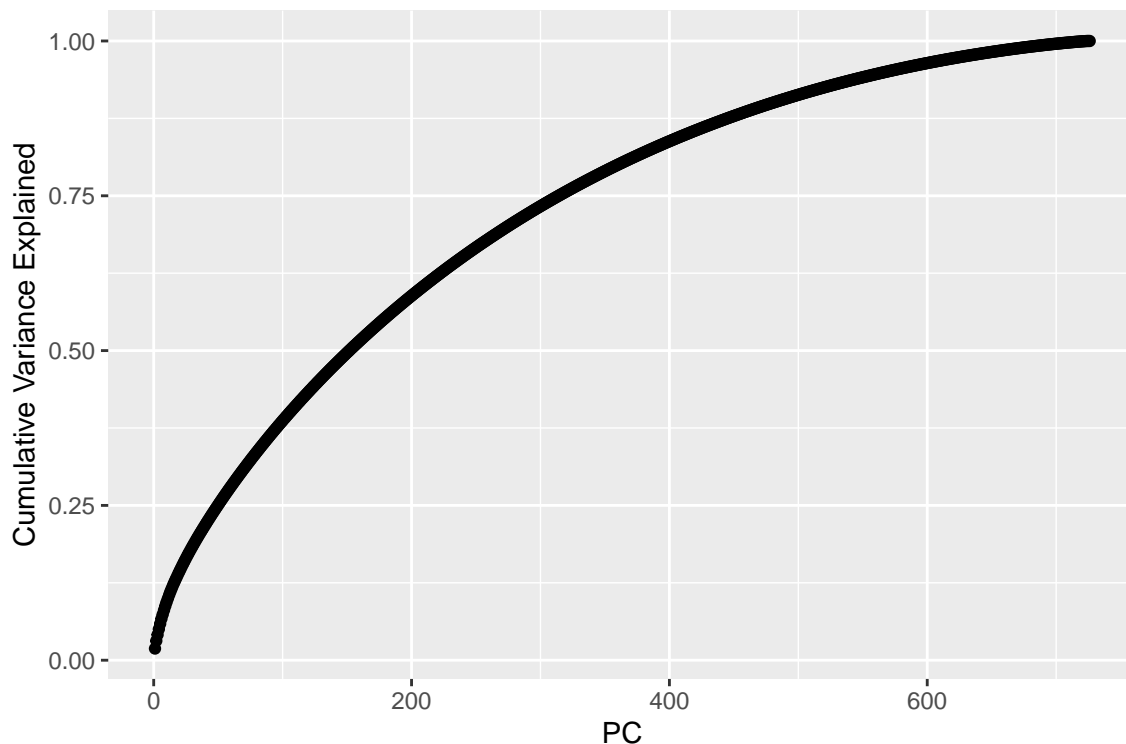
## <<DocumentTermMatrix (documents: 5, terms: 5)>>
## Non-/sparse entries: 6/19
## Sparsity          : 76%
## Maximal term length: 10
## Weighting         : term frequency (tf)
## Sample           :
##      Terms
## Docs aaron abandon able account accounting
##   1    1        1    1        1          3
##   2    0        0    0        0          0
##   3    0        0    0        0          0
##   4    0        0    0        0          0
##   5    1        0    0        0          0
```

We then filter words that are uninformative and appear in every document (i.e. all entries are 0).

Due to our filtering preference, we found that the words (features) between the training dataset and testing dataset matched up perfectly. Because of this, no other steps were needed to fix the features between the datasets.

Dimension Reduction

We then perform PCA on our training dataset. To ensure the principal components between both the training and testing dataset are equivalent, we project the observations of the test set into the PC space as defined by our training dataset. We choose the first 500 PCs to improve computational speed (~85% of variance explained).



Prediction

We use the random forest and k-Nearest Neighbor algorithms to determine a model for prediction.

We first try the random forest algorithm on our original dataset, rather than PC space. We have 725 words, and choose 70 as the number of variables to consider for each split as well as 500 trees for computational speed. We were able to achieve a classification accuracy of approximately 58%. *Note: we drop the feature (word) “next” from the dataset, as the random forest function throws an error when it encounters this word.*

```
## Warning: package 'randomForest' was built under R version 4.0.2
```

```
## randomForest 4.6-14
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
##
```

```
## Attaching package: 'randomForest'
```

```
## The following object is masked from 'package:ggplot2':
```

```
##
```

```
## margin
```

```
## The following object is masked from 'package:dplyr':
```

```
##
```

```
## combine
```

```
## [1] 0.5804
```

In random forest, we chose several hyperparameters that could balance between bias and variance, while trying to account for computational speed. We designated a subset of 20 variables for each split, and we grew 500 trees.

```
## [1] 0.4436
```

Using the random forest algorithm on our PC space, we were able to achieve a roughly 44% accuracy on our dataset. While it does reduce our dimensionality, it also reduces quite a bit of accuracy.

We then fit a k-Nearest Neighbors model on our training dataset, and chose to use $k = 6$.

```
## Warning: package 'class' was built under R version 4.0.2
```


```
##
```

```
## Attaching package: 'class'
```

```
## The following object is masked from 'package:reshape':
```

```
##
```

```
##      condense
```



In this situation, $k = 6$ provided the best accuracy.

In general, our random forest model fit on the original dataset (not PC space) performed the best.

Problem 6: Association Rule Mining

In this problem, we wanted to find association rules within a dataset containing baskets of grocery shoppers. We first take the necessary data pre-processing steps to read in the data, as well as format it in a way that allows the “arules” package in R to work. Below is a summary of the resulting object from the class “transactions” (also within the “arules” package).

In our dataset, we can see that whole milk is purchased in roughly 25% of the baskets.

```
## transactions as itemMatrix in sparse format with
## 9835 rows (elements/itemsets/transactions) and
## 169 columns (items) and a density of 0.02609146
##
## most frequent items:
##      whole milk other vegetables      rolls/buns      soda
##      2513      1903      1809      1715
##      yogurt      (Other)
##      1372      34055
##
## element (itemset/transaction) length distribution:
## sizes
##      1      2      3      4      5      6      7      8      9     10     11     12     13     14     15     16
## 2159 1643 1299 1005  855  645  545  438  350  246  182  117  78  77  55  46
##      17     18     19     20     21     22     23     24     26     27     28     29     32
##      29     14     14      9     11      4      6      1      1      1      1      3      1
##
##      Min. 1st Qu.  Median      Mean 3rd Qu.      Max.
##      1.000  2.000   3.000   4.409   6.000  32.000
##
## includes extended item information - examples:
##      labels
## 1 abrasive cleaner
## 2 artif. sweetener
## 3  baby cosmetics
```

We tested several thresholds for both confidence and lift. In general, what we were looking for were rules that had a high number of counts (in terms of the left-hand side of the rule) and high confidence as well as lift.

Below is a small sample of our baseline list of rules. We choose a confidence of 0.1 and a lift of 1. We choose a max length of 3 in order to find more general rules, instead of finding rules with larger lengths that may not occur as often. In total, there are 18597 rules.

```
##      lhs      rhs      support      confidence
## [1] {}      => {yogurt}      0.139501779 0.1395018
## [2] {}      => {rolls/buns}      0.183934926 0.1839349
## [3] {honey}  => {whole milk}      0.001118454 0.7333333
## [4] {soap}   => {whole milk}      0.001118454 0.4230769
## [5] {tidbits} => {soda}      0.001016777 0.4347826
## [6] {tidbits} => {rolls/buns}      0.001220132 0.5217391
## [7] {cocoa drinks} => {whole milk}      0.001321810 0.5909091
## [8] {snack products} => {soda}      0.001118454 0.3666667
## [9] {snack products} => {rolls/buns}      0.001118454 0.3666667
```

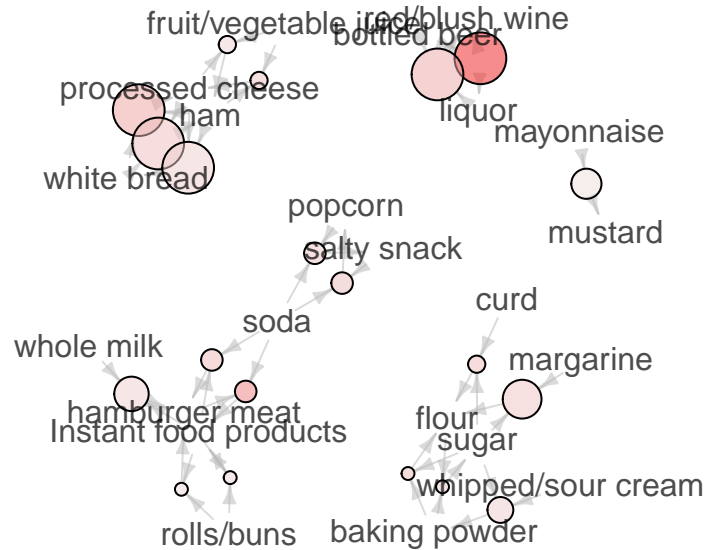
```

## [10] {pudding powder}    => {whole milk}      0.001321810 0.5652174
## [11] {cooking chocolate} => {whole milk}      0.001321810 0.5200000
## [12] {bathroom cleaner}  => {soda}             0.001016777 0.3703704
## [13] {bathroom cleaner}  => {other vegetables} 0.001016777 0.3703704
## [14] {nuts/prunes}       => {rolls/buns}      0.001016777 0.3030303
## [15] {nuts/prunes}       => {whole milk}      0.001220132 0.3636364
## [16] {brandy}           => {shopping bags}    0.001321810 0.3170732
## [17] {brandy}           => {rolls/buns}      0.001118454 0.2682927
## [18] {nut snack}         => {soda}             0.001016777 0.3225806
## [19] {potato products}   => {pastry}          0.001016777 0.3571429
## [20] {potato products}   => {whole milk}      0.001220132 0.4285714
##      coverage    lift    count
## [1] 1.000000000 1.000000 1372
## [2] 1.000000000 1.000000 1809
## [3] 0.001525165 2.870009   11
## [4] 0.002643620 1.655775   11
## [5] 0.002338587 2.493345   10
## [6] 0.002338587 2.836542   12
## [7] 0.002236909 2.312611   13
## [8] 0.003050330 2.102721   11
## [9] 0.003050330 1.993459   11
## [10] 0.002338587 2.212062   13
## [11] 0.002541942 2.035097   13
## [12] 0.002745297 2.123961   10
## [13] 0.002745297 1.914132   10
## [14] 0.003355363 1.647486   10
## [15] 0.003355363 1.423145   12
## [16] 0.004168785 3.218178   13
## [17] 0.004168785 1.458628   11
## [18] 0.003152008 1.849901   10
## [19] 0.002846975 4.014286   10
## [20] 0.002846975 1.677278   12

```

Graph for 20 rules

size: support (0.001 – 0.002)
color: lift (12.965 – 35.716)




In addition, from the graph, we can see by increasing the confidence (as well as reducing the minimum support required), we can see a drastic decrease in the number of rules in our list. These rules are generally good, because they have high count and have a good level of confidence.

```
## Apriori
##
## Parameter specification:
## confidence minval smax arem aval originalSupport maxtime support minlen
##          0.3   0.1   1 none FALSE                TRUE      5   0.05    1
## maxlen target  ext
##          3  rules TRUE
##
## Algorithmic control:
## filter tree heap memopt load sort verbose
##    0.1 TRUE TRUE  FALSE TRUE    2    TRUE
##
## Absolute minimum support count: 491
##
## set item appearances ...[0 item(s)] done [0.00s].
## set transactions ...[169 item(s), 9835 transaction(s)] done [0.00s].
## sorting and recoding items ... [28 item(s)] done [0.00s].
## creating transaction tree ... done [0.00s].
## checking subsets of size 1 2 done [0.00s].
## writing ... [3 rule(s)] done [0.00s].
## creating S4 object ... done [0.00s].

##      lhs                rhs      support   confidence coverage  lift
## [1] {yogurt}            => {whole milk} 0.05602440 0.4016035 0.1395018 1.571735
```

```
## [2] {rolls/buns}      => {whole milk} 0.05663447 0.3079049 0.1839349 1.205032
## [3] {other vegetables} => {whole milk} 0.07483477 0.3867578 0.1934926 1.513634
##      count
## [1] 551
## [2] 557
## [3] 736
```



By decreasing the required minimum support again, we can achieve a slightly larger list (size 37) of rules. However, in general, the counts for each rule is low (around 200 each rule).

Using the same rules, we can filter out the items that have a high value for lift. In general, since lift is high, it implies that the items on either side of the rule are not independent from one another (i.e. they are complementary). An interesting rule to note on this list is the second rule - the count is high and the confidence is good. It also makes sense - typically, when a shopper is buying vegetables, they probably won't just stop at root vegetables.

```
## Apriori
##
## Parameter specification:
## confidence minval smax arem aval originalSupport maxtime support minlen
##          0.3    0.1    1 none FALSE                TRUE      5    0.02    1
## maxlen target  ext
##          3  rules TRUE
##
## Algorithmic control:
## filter tree heap memopt load sort verbose
##      0.1 TRUE TRUE  FALSE TRUE    2    TRUE
##
```



```

## Absolute minimum support count: 196
##
## set item appearances ...[0 item(s)] done [0.00s].
## set transactions ...[169 item(s), 9835 transaction(s)] done [0.00s].
## sorting and recoding items ... [59 item(s)] done [0.00s].
## creating transaction tree ... done [0.00s].
## checking subsets of size 1 2 3

## Warning in apriori(baskets.trans, parameter = list(support = 0.02, confidence
## = 0.3, : Mining stopped (maxlen reached). Only patterns up to a length of 3
## returned!


## done [0.00s].
## writing ... [37 rule(s)] done [0.00s].
## creating S4 object ... done [0.00s].

##      lhs                                rhs          support    confidence
## [1] {whipped/sour cream}                => {other vegetables} 0.02887646 0.4028369
## [2] {root vegetables}                   => {other vegetables} 0.04738180 0.4347015
## [3] {root vegetables,whole milk}        => {other vegetables} 0.02318251 0.4740125
## [4] {other vegetables,whole milk}       => {root vegetables} 0.02318251 0.3097826
## [5] {other vegetables,yogurt}           => {whole milk}       0.02226741 0.5128806
## [6] {whole milk,yogurt}                 => {other vegetables} 0.02226741 0.3974592
##      coverage lift    count
## [1] 0.07168277 2.081924 284
## [2] 0.10899847 2.246605 466
## [3] 0.04890696 2.449770 228
## [4] 0.07483477 2.842082 228
## [5] 0.04341637 2.007235 219
## [6] 0.05602440 2.054131 219

```

Graph for 6 rules

size: support (0.022 – 0.047)
color: lift (2.007 – 2.842)



We then reduced the minimum required support even further and increased the maximum length of each rule. In general, most of these rules are good (if we disregard the low count). They all have lifts greater than 2, which is indicative of their possible non-independence.

```
## Apriori
##
## Parameter specification:
## confidence minval smax arem aval originalSupport maxtime support minlen
##          0.3   0.1   1 none FALSE                TRUE     5   0.01     1
## maxlen target  ext
##          4  rules TRUE
##
## Algorithmic control:
## filter tree heap memopt load sort verbose
##    0.1 TRUE TRUE  FALSE TRUE    2    TRUE
##
## Absolute minimum support count: 98
##
## set item appearances ...[0 item(s)] done [0.00s].
## set transactions ...[169 item(s), 9835 transaction(s)] done [0.00s].
## sorting and recoding items ... [88 item(s)] done [0.00s].
## creating transaction tree ... done [0.00s].
## checking subsets of size 1 2 3 4

## Warning in apriori(baskets.trans, parameter = list(support = 0.01, confidence
## = 0.3, : Mining stopped (maxlen reached). Only patterns up to a length of 4
## returned!
```


```
## done [0.00s].
## writing ... [125 rule(s)] done [0.00s].
## creating S4 object ... done [0.00s].
```

##	lhs	rhs	support	confidence	coverage	lift	count
## [1]	{onions}	=> {other vegetables}	0.01423488	0.4590164	0.03101169	2.372268	140
## [2]	{berries}	=> {yogurt}	0.01057448	0.3180428	0.03324860	2.279848	104
## [3]	{hamburger meat}	=> {other vegetables}	0.01382816	0.4159021	0.03324860	2.149447	136
## [4]	{cream cheese }	=> {yogurt}	0.01240468	0.3128205	0.03965430	2.242412	122
## [5]	{chicken}	=> {other vegetables}	0.01789527	0.4170616	0.04290798	2.155439	176
## [6]	{beef}	=> {root vegetables}	0.01738688	0.3313953	0.05246568	3.040367	171
## [7]	{curd}	=> {yogurt}	0.01728521	0.3244275	0.05327911	2.325615	170
## [8]	{whipped/sour cream}	=> {other vegetables}	0.02887646	0.4028369	0.07168277	2.081924	284
## [9]	{root vegetables}	=> {other vegetables}	0.04738180	0.4347015	0.10899847	2.246605	466
## [10]	{curd, yogurt}	=> {whole milk}	0.01006609	0.5823529	0.01728521	2.279125	99
## [11]	{curd, whole milk}	=> {yogurt}	0.01006609	0.3852140	0.02613116	2.761356	99
## [12]	{pork, whole milk}	=> {other vegetables}	0.01016777	0.4587156	0.02216573	2.370714	100
## [13]	{butter, other vegetables}	=> {whole milk}	0.01148958	0.5736041	0.02003050	2.244885	113
## [14]	{butter, whole milk}	=> {other vegetables}	0.01148958	0.4169742	0.02755465	2.154987	113
## [15]	{domestic eggs, other vegetables}	=> {whole milk}	0.01230300	0.5525114	0.02226741	2.162336	121
## [16]	{domestic eggs, whole milk}	=> {other vegetables}	0.01230300	0.4101695	0.02999492	2.119820	121
## [17]	{fruit/vegetable juice, whole milk}	=> {other vegetables}	0.01047280	0.3931298	0.02663955	2.031756	103
## [18]	{whipped/sour cream, yogurt}	=> {other vegetables}	0.01016777	0.4901961	0.02074225	2.533410	100
## [19]	{other vegetables, whipped/sour cream}	=> {yogurt}	0.01016777	0.3521127	0.02887646	2.524073	100
## [20]	{whipped/sour cream, yogurt}	=> {whole milk}	0.01087951	0.5245098	0.02074225	2.052747	107
## [21]	{whipped/sour cream, whole milk}	=> {yogurt}	0.01087951	0.3375394	0.03223183	2.419607	107
## [22]	{whipped/sour cream, whole milk}	=> {other vegetables}	0.01464159	0.4542587	0.03223183	2.347679	144
## [23]	{other vegetables, pip fruit}	=> {whole milk}	0.01352313	0.5175097	0.02613116	2.025351	133
## [24]	{pip fruit, whole milk}	=> {other vegetables}	0.01352313	0.4493243	0.03009659	2.322178	133
## [25]	{citrus fruit, root vegetables}	=> {other vegetables}	0.01037112	0.5862069	0.01769192	3.029608	102
## [26]	{citrus fruit, other vegetables}	=> {root vegetables}	0.01037112	0.3591549	0.02887646	3.295045	102
## [27]	{citrus fruit, whole milk}	=> {yogurt}	0.01026945	0.3366667	0.03050330	2.413350	101
## [28]	{citrus fruit, whole milk}	=> {other vegetables}	0.01301474	0.4266667	0.03050330	2.205080	128
## [29]	{root vegetables, tropical fruit}	=> {other vegetables}	0.01230300	0.5845411	0.02104728	3.020999	121

## [30] {other vegetables, ## tropical fruit}	=> {root vegetables}	0.01230300	0.3427762	0.03589222	3.144780	121
## [31] {root vegetables, ## tropical fruit}	=> {whole milk}	0.01199797	0.5700483	0.02104728	2.230969	118
## [32] {tropical fruit, ## yogurt}	=> {other vegetables}	0.01230300	0.4201389	0.02928317	2.171343	121
## [33] {other vegetables, ## tropical fruit}	=> {yogurt}	0.01230300	0.3427762	0.03589222	2.457146	121
## [34] {tropical fruit, ## yogurt}	=> {whole milk}	0.01514997	0.5173611	0.02928317	2.024770	149
## [35] {tropical fruit, ## whole milk}	=> {yogurt}	0.01514997	0.3581731	0.04229792	2.567516	149
## [36] {tropical fruit, ## whole milk}	=> {other vegetables}	0.01708185	0.4038462	0.04229792	2.087140	168
## [37] {root vegetables, ## yogurt}	=> {other vegetables}	0.01291307	0.5000000	0.02582613	2.584078	127
## [38] {root vegetables, ## yogurt}	=> {whole milk}	0.01453991	0.5629921	0.02582613	2.203354	143
## [39] {rolls/buns, ## root vegetables}	=> {other vegetables}	0.01220132	0.5020921	0.02430097	2.594890	120
## [40] {rolls/buns, ## root vegetables}	=> {whole milk}	0.01270971	0.5230126	0.02430097	2.046888	125
## [41] {root vegetables, ## whole milk}	=> {other vegetables}	0.02318251	0.4740125	0.04890696	2.449770	228
## [42] {other vegetables, ## whole milk}	=> {root vegetables}	0.02318251	0.3097826	0.07483477	2.842082	228
## [43] {other vegetables, ## yogurt}	=> {whole milk}	0.02226741	0.5128806	0.04341637	2.007235	219
## [44] {whole milk, ## yogurt}	=> {other vegetables}	0.02226741	0.3974592	0.05602440	2.054131	219

Graph for 20 rules

size: support (0.01 – 0.023)
color: lift (2.326 – 3.295)



We then looked for substitute items. After reducing support and maximum length slightly, we found an interesting set of rules that depicted this. In general, we believe that it's safe to assume that most people who go shopping are not there to buy shopping bags, but rather are there to buy whole milk.

```
## Apriori
##
## Parameter specification:
## confidence minval smax arem aval originalSupport maxtime support minlen
##          0.01   0.1   1 none FALSE                TRUE      5  0.005     1
## maxlen target  ext
##          3  rules TRUE
##
## Algorithmic control:
## filter tree heap memopt load sort verbose
##    0.1 TRUE TRUE  FALSE TRUE    2    TRUE
##
## Absolute minimum support count: 49
##
## set item appearances ...[0 item(s)] done [0.00s].
## set transactions ...[169 item(s), 9835 transaction(s)] done [0.00s].
## sorting and recoding items ... [120 item(s)] done [0.00s].
## creating transaction tree ... done [0.00s].
## checking subsets of size 1 2 3

## Warning in apriori(baskets.trans, parameter = list(support = 0.005, confidence
## = 0.01, : Mining stopped (maxlen reached). Only patterns up to a length of 3
## returned!
```

```
## done [0.00s].
## writing ... [2090 rule(s)] done [0.00s].
## creating S4 object ... done [0.00s].

##      lhs                rhs          support    confidence coverage
## [1] {shopping bags} => {whole milk}  0.02450432 0.24871001 0.09852567
## [2] {whole milk}    => {shopping bags} 0.02450432 0.09590131 0.25551601
##      lift      count
## [1] 0.9733637 241
## [2] 0.9733637 241
```

In general, most of the rules make sense. Root vegetables aren't going to be the only vegetables you buy, you'll probably buy other vegetables. If you're getting yogurt, you're going to be in the dairy section, so you'll probably pick up some milk too. And finally, if you're going to make a sandwich, you'll probably need both white bread and ham.